


COMS W4733: Computational Aspects of Robotics

Lecture 11: Bug Algorithms 1



Slide materials from H. Choset, G. D. Hager, and Z. Dodds


Instructor: Tony Dear

Planning Algorithms

- Up to now: Continuous trajectories represented as mathematical functions
 - Easy to formulate, guaranteed solution given initial and final conditions
 - But *extremely* limited to obstacle-free, static, fully observable environments!
-
- Robots operate in the real world
 - Full of obstacles (possibly dynamic), oftentimes not fully observable
 - Plans and trajectories need to update in response to environment feedback
-
- Choice of algorithms will depend on what we know and what we don't


Bug Algorithms

- How do bugs traverse the environment?
- Assume only local knowledge of surroundings
 - Tactile sensing
 - Finite distance sensing
- Behaviors are simple and can be enumerated
 - Follow a wall
 - Move toward a goal
- Overall planning is *incremental* and *reactive*



Environment Setup


- Assume point robot—not worrying about kinematics or constraints here!
- Environment is bounded, finite number of obstacles
 - Robot knows its global position
 - Robot is able to measure distance between points
 - Does not know layout of obstacles
- Workspace: (x, y) or $(x, y, \theta) \in W$
- Set of obstacles: WO_i
- Free workspace: $W_{free} = W - \cup_i WO_i$



• Start


Bug 0


- If no obstacles, best path to the goal is a simple straight line
- May encounter obstacles between q_{start} and q_{goal}
 - *Hit point* on obstacle i : q_i^H
 - *Leave point* on obstacle i : q_i^L
- *Path* can be represented as a sequence of hit-leave point pairs
- While not at goal:
 - If at obstacle according to sensors:
 - Follow obstacle until we can head toward goal
 - Else: Head toward goal



Bug 0


- When does bug 0 fail?
- Drawback: Bug 0 has no memory
- Going around in loops can be a problem
- Can we improve Bug 0 without adding any features?
- What if we always turn right rather than left?
- Unfortunately we wouldn't know when to switch






Bug 1


- $q_0^L \leftarrow q_{start}, i = 1$
- Loop:
 - Repeat: From q_{i-1}^L move toward q_{goal}
 - If goal reached: Exit and return success
 - If obstacle WO_i encountered at q_i^H :
 - Circumnavigate WO_i and record closest
 - If goal reached: Return success
 - If q_i^H re-encountered: return to q_i^L
 - If move toward q_{goal} encounters obstacle:
 - Return failure
 - Else: $i \leftarrow i + 1$



Bug 1 Examples



Bug 1 Examples




Bug 1 Completeness

- Suppose Bug 1 never terminates...
 - But each leave point is closer to the goal than the hit point, since we pick q_i^L to be the closest point on the obstacle: $d(q_i^L, q_{goal}) < d(q_i^H, q_{goal})$
 - Each hit point is also closer to the goal than the last hit point, since the robot moved toward the obstacle from q_{i-1}^L to arrive at q_i^H
 - There a finite number of hit/leave point pairs, which the robot will eventually exhaust
- Suppose Bug 1 incorrectly returns failure...
 - Then it will attempt to leave from the closest leave point q_i^L that will run into an obstacle
 - But there must be at least one other possible leave point q_i^{L*} on the obstacle along line to goal
 - Since a path exists, robot would have encountered q_i^{L*} while circumnavigating obstacle and would not leave from q_i^L

Bug 1 Performance


- What are the best- and worst-case scenarios?
- Best case: Robot goes straight to goal
- Distance = D (distance from start to goal)
- Worst case: Robot travels D distance, plus completely circumnavigating all obstacles
- Distance = $D + 1.5 \sum_i P_i$
- P_i = perimeter of WO_i



Bug 2

- Do we really need to explore every obstacle's perimeter all the time?
- We already had some idea of the “shortest path” to the goal from the start!
- New idea: Just follow the **m-line** whenever we are able to instead of computing new shortest paths
- Only need to remember m-line and hit points


Call the line from the starting point to the goal the **m-line**



Bug 2


- $q_0^L \leftarrow q_{start}, i = 1$
- Loop:
 - Repeat: From q_{i-1}^L move toward q_{goal} along m-line
 - If goal reached: Exit and return success
 - If obstacle WO_i encountered at q_i^H :
 - Repeat: Turn left or right and follow boundary of WO_i
 - If goal reached: Return success
 - If q_i^H re-encountered: Return failure
 - If m-line encountered at m s.t. $d(m, q_{goal}) < d(q_i^H, q_{goal})$ and line(m, q_{goal}) does not encounter obstacle:
 - Stop following WO_i
 - $q_i^L \leftarrow m; i \leftarrow i + 1$

Call the line from the starting point to the goal the **m-line**




Bug 2


- Why do we need the distance condition when circling WO_i ?
- If m-line encountered at m s.t. $d(m, q_{goal}) < d(q_i^H, q_{goal})$ and line(m, q_{goal}) does not encounter obstacle
- Is it possible to re-encounter m-line such that we're actually farther away from the goal than when we started?
- In other words, can we somehow follow an obstacle to move behind the start or hit point?
- Distance condition ensures we don't get stuck in infinite loop!



Bug 2 Examples




Bug 2 Examples




Bug 1 vs Bug 2

Bug 2 beats Bug 1?




Bug 1 beats Bug 2?



Bug 2 Performance

- Best case: Robot goes straight to goal, distance = D
- Same as Bug 1
- What is the worst case?
- Each time we follow the m-line and find a hit point q_i^H , we could potentially traverse the perimeter!
- Suppose n_i = number of intersections with WO_i
- Worst case: distance = $D + \frac{1}{2} \sum_i n_i P_i$



Bug Algorithm Properties

- Bug 1 is an exhaustive search algorithm
- Circumnavigate every obstacle, check all options before choosing the best one
- Bug 2 is a greedy search algorithm
- Takes the first thing that looks good—leave points along stored m-line
- Bug 2 can outperform Bug 1 in many situations, but not always
- Bug 1 has more predictable performance and results
- Both algorithms are *complete*—either return a path or failure if none exists