


COMS W4733: Computational Aspects of Robotics

Lecture 22: Graph SLAM



Instructor: Tony Dear

Slide materials from P. Allen, M. Bennewitz, W. Burgard, D. Fox, M. Kaess, L. Spinello, C. Stachniss, D. Tipaldi, et al.

State Estimation

- Belief distribution

$$B(\mathbf{x}_t) = p(\mathbf{x}_t \mid \mathbf{u}_{1:t}, \mathbf{z}_{1:t})$$

Robot state at time t

Actions from 1 to t


Observations from 1 to t

- Transition model

$$p(\mathbf{x}_t \mid \mathbf{x}_{t-1}, \mathbf{u}_t)$$

- Observation model

$$p(\mathbf{z}_t \mid \mathbf{x}_t)$$



Bayes Filter Algorithm

Algorithm **Bayes_filter**($B(x)$, d):

1. $\eta = 0$
2. if d is an *action* data item u then
3. for all x do
4. $\bar{B}(x) = \int p(x|x', u)B(x')dx'$
5. if d is a *perceptual* data item z then
6. for all x do
7. $\bar{B}(x) = p(z|x)\bar{B}(x)$
8. $\eta = \eta + \bar{B}(x)$
9. for all x do
10. $\bar{B}(x) = \eta^{-1}\bar{B}(x)$
11. return $\bar{B}(x)$

Prediction:

$$B'(x_t) = \int \underbrace{p(x_t|x_{t-1}, u_t)}_{\text{Transition model}} B(x_{t-1}) dx_{t-1}$$

Observation:

$$B(x_t) = \eta^{-1} \underbrace{p(z_t|x_t)}_{\text{Observation model}} B'(x_t)$$

Kalman Filter

- Start with current belief distribution:

$$Bel(\mathbf{x}_{k-1}) \sim N(\hat{\mathbf{x}}_{k-1|k-1}, \mathbf{P}_{k-1|k-1})$$

- Predict according to transition model:

$$\begin{aligned}\hat{\mathbf{x}}_{k|k-1} &= \mathbf{F}_k \hat{\mathbf{x}}_{k-1|k-1} + \mathbf{B}_k \mathbf{u}_k \\ \mathbf{P}_{k|k-1} &= \mathbf{F}_k \mathbf{P}_{k-1|k-1} \mathbf{F}_k^T + \mathbf{Q}_k\end{aligned}$$


- Update according to observation model and measurement:

$$\mathbf{K}_k = \mathbf{P}_{k|k-1} \mathbf{H}_k^T (\mathbf{H}_k \mathbf{P}_{k|k-1} \mathbf{H}_k^T + \mathbf{R}_k)^{-1}$$

$$\hat{\mathbf{x}}_{k|k} = \hat{\mathbf{x}}_{k|k-1} + \mathbf{K}_k (\mathbf{z}_k - \mathbf{H}_k \hat{\mathbf{x}}_{k|k-1})$$

$$\mathbf{P}_{k|k} = \mathbf{P}_{k|k-1} - \mathbf{K}_k \mathbf{H}_k \mathbf{P}_{k|k-1}$$

Example: Kalman Filter




Particle Filter

Algorithm **Particle_filter**($X_{t-1}, \mathbf{u}_t, \mathbf{z}_t$):

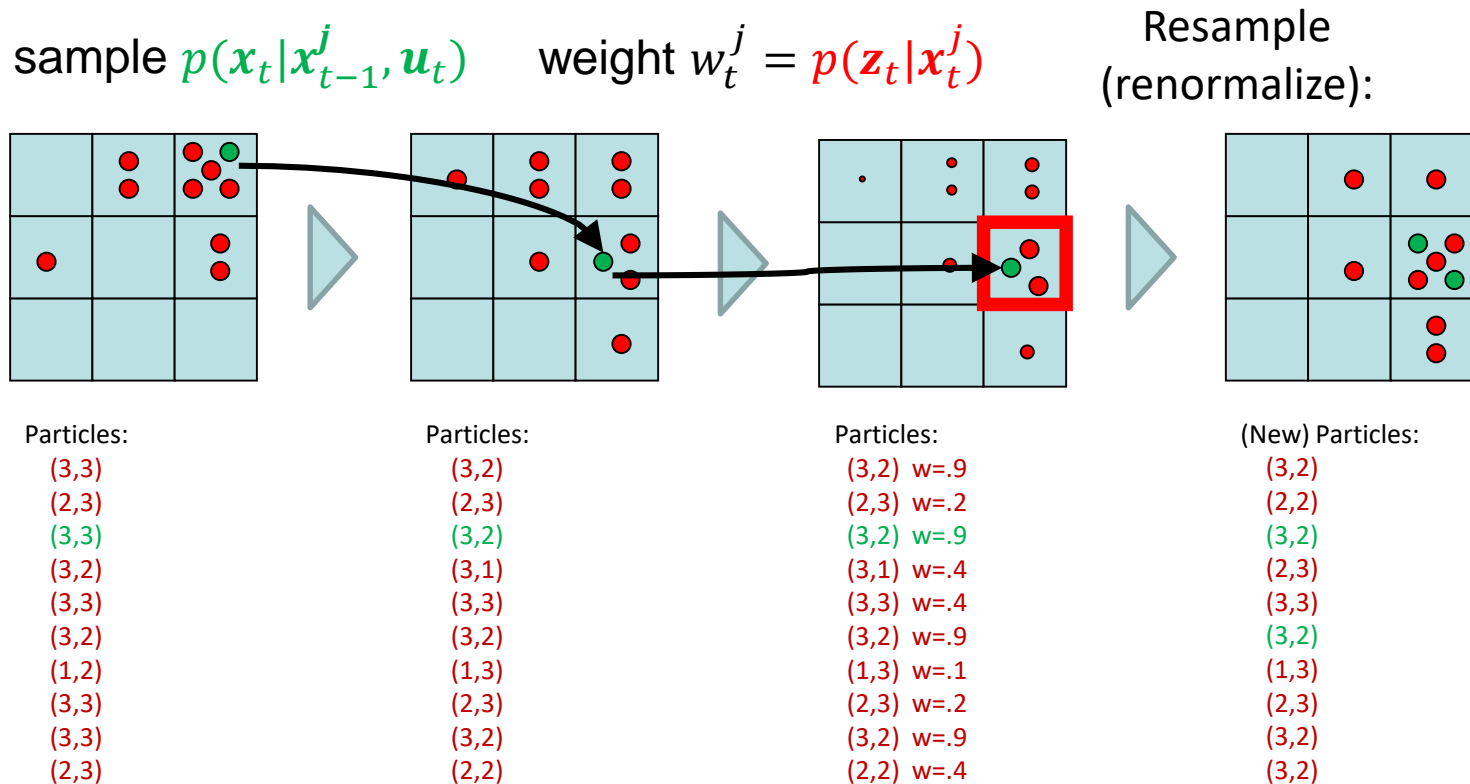
1. $\bar{X}_t = \emptyset, X_t = \emptyset$
2. **for** each particle x_{t-1}^j in X_{t-1} **do**
3. sample x_t^j from $p(x_t | x_{t-1}^j, \mathbf{u}_t)$
4. compute weight $w_t^j = p(\mathbf{z}_t | x_t^j)$
5. insert (x_t^j, w_t^j) into \bar{X}_t
6. **for all** j **do**
7. sample $i \in \{1, 2, \dots, J\}$ with prob $\frac{w_t^j}{\sum w_t^j}$
8. insert x_t^i from \bar{X}_t into X_t
9. **return** X_t

$$B'(\mathbf{x}_t) = \int p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{u}_t) B(\mathbf{x}_{t-1}) d\mathbf{x}_{t-1}$$


$$B(\mathbf{x}_t) = \eta^{-1} p(\mathbf{z}_t | \mathbf{x}_t) B'(\mathbf{x}_t)$$



Example: Particle Filter



Example: Particle Filter



Example: Particle Filter



Navigation Tasks


- **Localization:** Given a map, infer the robot's location
- Variations: Known/unknown initial location, false location (kidnapped robot)
- **Mapping:** Given location(s), infer a *map* of the environment
 - Manually driven or fully autonomous?
 - How to define exploration of the area?
 - Termination condition?



Map Representations


Grid-based:

Collection of discretized
obstacle / free-space pixels




Feature-based:

Collection of landmark
locations and correlated
uncertainties



Topological:

Collection of nodes and
interconnections




Map Representations

	Grid-Based	Feature-Based	Topological
Resolution vs. Scale	Discrete localization	Arbitrary localization	Localize to nodes
Computational Complexity	Grid size <i>and</i> resolution	Landmark covariance (N^2)	Minimal complexity
Obstacle Avoidance	Discretized obstacles	Only structured obstacles	GVG defines the safest path
Exploration Strategies	Frontier-based exploration	No inherent exploration	Graph exploration


SLAM

- **Simultaneous localization and mapping (SLAM)** entails learning a map and locating the robot at the same time!
- Clearly harder than either problem separately
- Main approach: Add the map (e.g., in the form of feature locations) to the state and update using Bayesian filtering
- Map is another hidden state!



Feature-Based SLAM

- At any given time, robot may be able to observe some number of landmarks
- These measurements are always *relative*, not absolute
- Errors in robot pose estimate and map representation are correlated!



Bayes Filter with Landmarks


- Belief distribution: $B(\mathbf{x}_t, \mathbf{m}) = p(\mathbf{x}_t, \mathbf{m} | \mathbf{u}_{1:t}, \mathbf{z}_{1:t})$

- Prediction:

$$B'(\mathbf{x}_t, \mathbf{m}) = \int \underbrace{p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{u}_t)}_{\text{Transition model}} B(\mathbf{x}_{t-1}, \mathbf{m}) d\mathbf{x}_{t-1}$$

- Observation:

$$B(\mathbf{x}_t, \mathbf{m}) = \eta^{-1} \underbrace{p(\mathbf{z}_t | \mathbf{x}_t, \mathbf{m})}_{\text{Observation model}} B'(\mathbf{x}_t, \mathbf{m})$$




Graph SLAM

- We need to be able to correlate different observations of the same landmarks
- Idea: Use a graph to capture structure and *constraints* between observations
- Nodes: Poses of the robot and landmark locations
- Edges: Spatial constraints between poses due to successive transitions or between poses and landmarks due to observations
- Once graph is constructed, use optimization to minimize the total *error* as the difference between predicted and actual observations




Graph Construction

- n nodes \mathbf{x}_i , one for each robot pose at time t_i


- *Odometry* edge between \mathbf{x}_i and \mathbf{x}_{i+1}



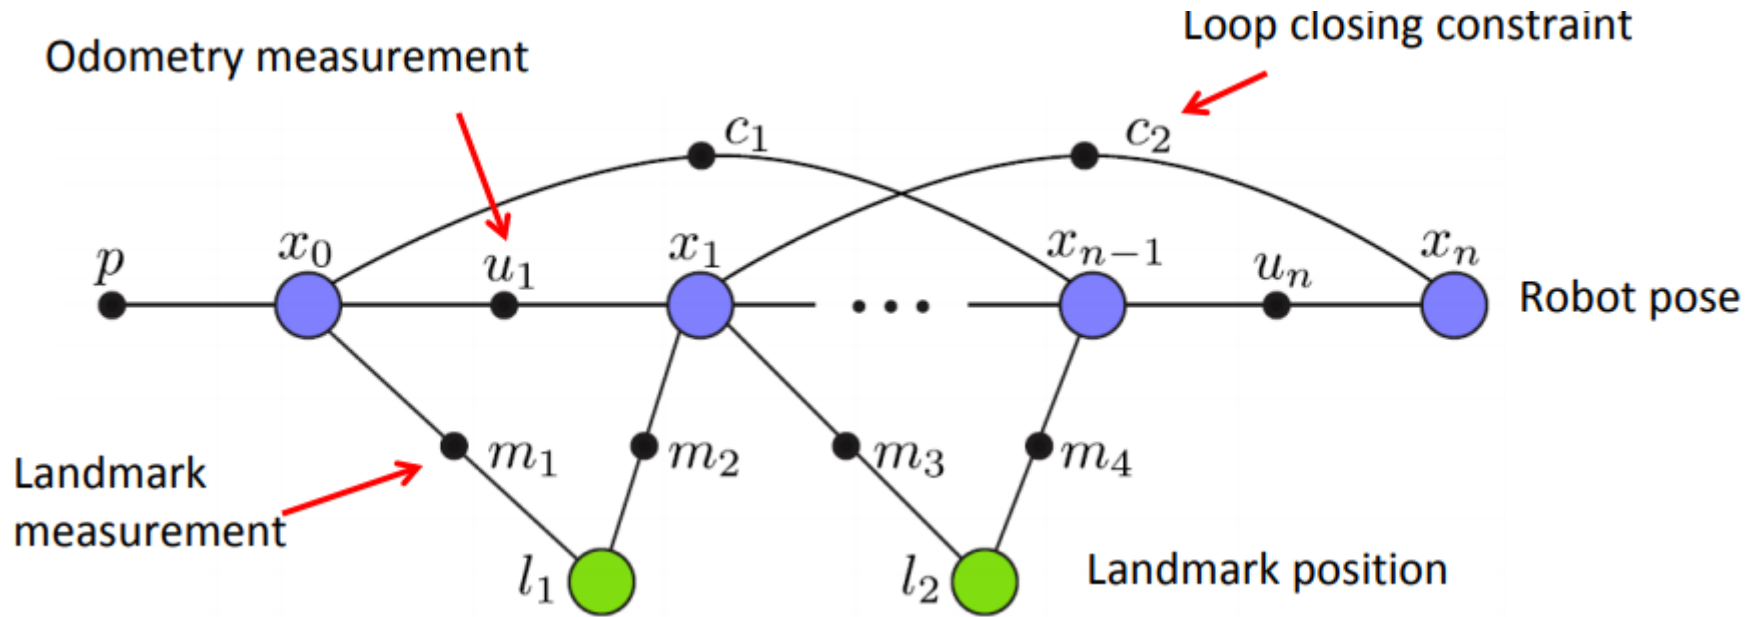
- *Observation* edge if observed \mathbf{l}_j from \mathbf{x}_i



Graph Construction



Graph Construction



Linear Optimization


- If our models are linear (or can be linearized), then we can solve for best-fitting poses over time and landmark locations using least-squares

- Ex: 1-D state, 3 timesteps, 1 landmark

$$AX = b$$

Can solve for X by
“pseudo-inverting” A


$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 1 & -1 & 0 & 0 \\ 0 & 1 & -1 & 0 \\ 1 & 0 & 0 & -1 \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & -1 \end{bmatrix} \begin{bmatrix} X_0 \\ X_1 \\ X_2 \\ L \end{bmatrix} = \begin{bmatrix} -3 \\ -5 \\ -3 \\ -10 \\ -5 \\ -2 \end{bmatrix}$$



Adding Uncertainty

- Recall form of left pseudo-inverse: $X = (A^T A)^{-1} A^T b$
- How to deal with uncertainty from transitions or observations?
- They will *weight* our constraints—the greater the uncertainty (variance), the less that we should trust the particular constraint
- W is block-diagonal, containing 1/variance terms:

$$X = (A^T W A)^{-1} A^T W b$$




$$W = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 5 \end{bmatrix}$$


Full Graph SLAM Problem

- In the most general case, each constraint represents an *error*
- Implicit assumption that this is zero-mean, normally distributed
- Odometry constraints weighted by \mathbf{Q}_t^{-1}
- Observation constraints weighted by \mathbf{R}_t^{-1}
- May also have covariance for initial state \mathbf{x}_0
- Loss function to be minimized (may be linearized):



$$J = \mathbf{x}_0^T \Omega_0 \mathbf{x}_0 + \sum_t (\mathbf{x}_t - f(\mathbf{x}_{t-1}, u_t))^T \mathbf{R}_t^{-1} (\mathbf{x}_t - f(\mathbf{x}_{t-1}, u_t)) + \sum_{t,i} (\mathbf{z}_t^i - h(\mathbf{x}_t, m_i))^T \mathbf{Q}_t^{-1} (\mathbf{z}_t^i - h(\mathbf{x}_t, m_i))$$



Full Graph SLAM Problem



Example: Graph SLAM



KUKA Halle 22, courtesy P. Pfaff

Summary

- Graph SLAM addresses full problem by optimizing (sparse) link graph between poses and landmarks
- Inference is baked into costs on edges
- Maximum likelihood estimate of both poses and map
- Still haven't really addressed the data association problem
- Long trajectories require more storage; method better suited for offline use
- Next time: Back to probabilistic methods (Kalman and particle filters!)