


COMS W4733: Computational Aspects of Robotics

Lecture 23: EKF-SLAM and FastSLAM




Instructor: Tony Dear

Slide materials from K. Arras, W. Burgard, J. Castellanos, G. Grisetti, J. Leonard, M. Montemerlo, C. Stachniss, et al.

SLAM

- **Simultaneous localization and mapping (SLAM)** entails learning a map and locating the robot at the same time!
- Clearly harder than either problem separately
- Main approach: Add the map (e.g., in the form of feature locations) to the state and update using Bayesian filtering
- Map is another hidden state!



Bayes Filter with Landmarks


- Belief distribution: $B(\mathbf{x}_t, \mathbf{m}) = p(\mathbf{x}_t, \mathbf{m} | \mathbf{u}_{1:t}, \mathbf{z}_{1:t})$

- Prediction:

$$B'(\mathbf{x}_t, \mathbf{m}) = \int \underbrace{p(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{u}_t)}_{\text{Transition model}} B(\mathbf{x}_{t-1}, \mathbf{m}) d\mathbf{x}_{t-1}$$

- Observation:

$$B(\mathbf{x}_t, \mathbf{m}) = \eta^{-1} \underbrace{p(\mathbf{z}_t | \mathbf{x}_t, \mathbf{m})}_{\text{Observation model}} B'(\mathbf{x}_t, \mathbf{m})$$




EKF for SLAM

- As with vanilla state estimation, we assume that our belief over the augmented state with landmark locations is Gaussian
- Run through standard Kalman operations to update robot state and map


$$\mu = \begin{pmatrix} x \\ y \\ \theta \\ m_{1x} \\ m_{1y} \\ \vdots \\ m_{nx} \\ m_{ny} \end{pmatrix} = \begin{pmatrix} X \\ m_1 \\ \vdots \\ m_n \end{pmatrix} \quad \Sigma = \begin{pmatrix} \Sigma_{XX} & \Sigma_{Xm_1} & \cdots & \Sigma_{Xm_n} \\ \Sigma_{m_1X} & \Sigma_{m_1m_1} & \cdots & \Sigma_{m_1m_n} \\ \vdots & \vdots & \ddots & \vdots \\ \Sigma_{m_nX} & \Sigma_{m_nm_1} & \cdots & \Sigma_{m_nm_n} \end{pmatrix}$$

EKF SLAM Steps


1. State transition




2. Predict measurement




3. Make measurement




4. Associate data




5. Update belief




Example: SLAM Without Landmarks




Example: SLAM With Landmarks



Example: SLAM With Landmarks




Example: SLAM With Landmarks




Robot also re-observes
the first two features,
resulting in uncertainty
reduction for ALL state
components

Example: SLAM With Landmarks



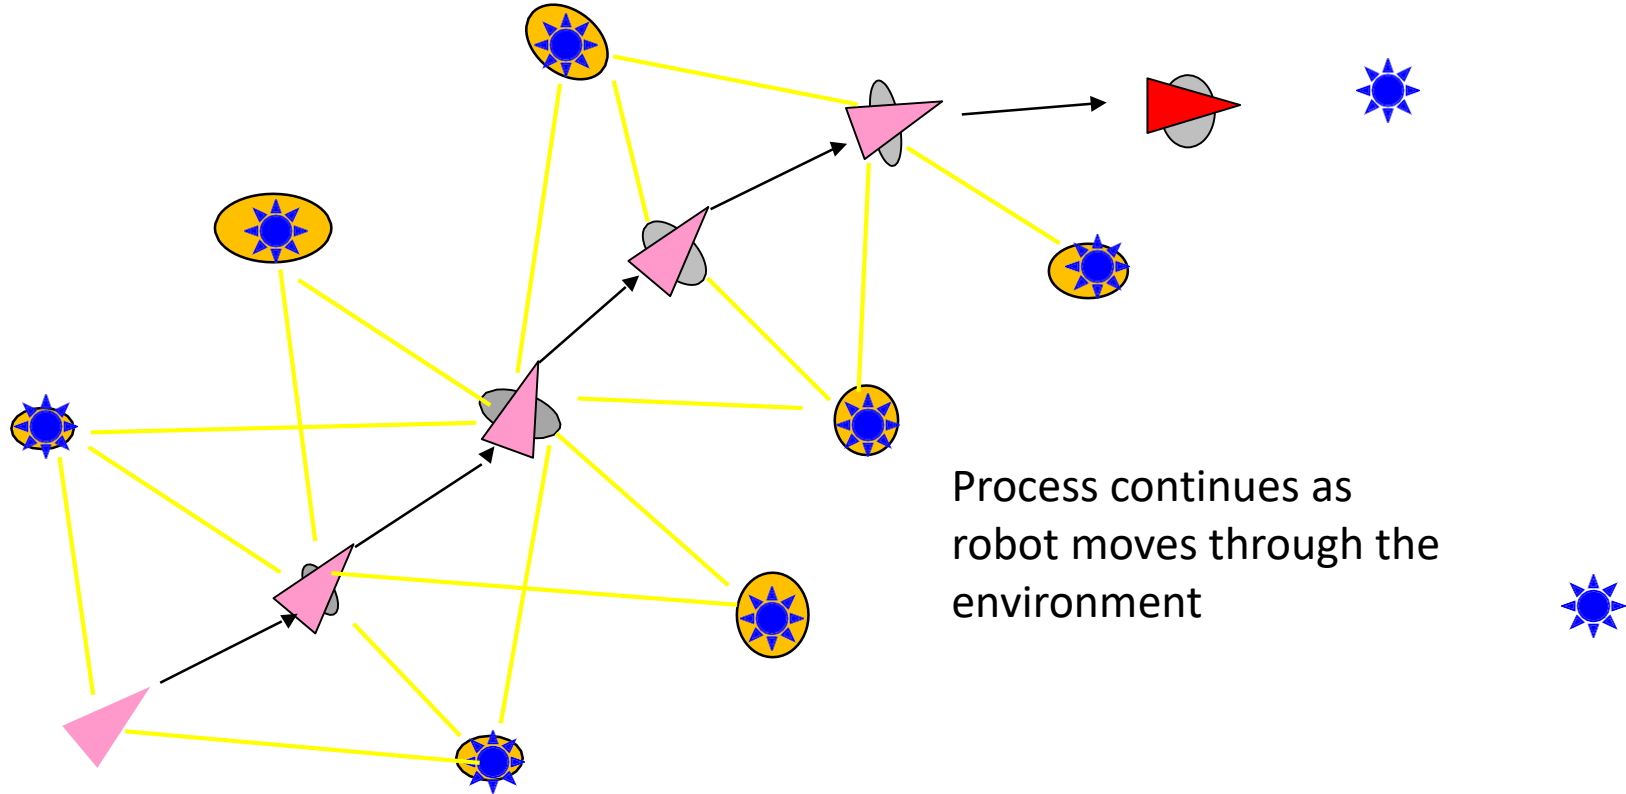
Robot observes another two new landmarks and adds them to the state

Example: SLAM With Landmarks



Robot also re-observes
the first four features,
resulting in uncertainty
reduction for ALL state
components

Example: SLAM With Landmarks



Transition Update: Mean

$$\hat{\mathbf{x}}_{k|k-1} = \mathbf{f}(\hat{\mathbf{x}}_{k-1|k-1}, \mathbf{u}_k)$$

- State transition utilizes the full, possibly nonlinear, system model \mathbf{f}
- Map landmark predictions should remain unchanged if stationary

- Ex:

$$\mathbf{f}(\hat{\mathbf{x}}_{k-1|k-1}, \mathbf{u}_k) = \begin{pmatrix} \hat{x} - \frac{v_k}{\omega_k} \sin \hat{\theta} + \frac{v_k}{\omega_k} \sin(\hat{\theta} + \omega_k \Delta t) \\ \hat{y} + \frac{v_k}{\omega_k} \cos \hat{\theta} - \frac{v_k}{\omega_k} \cos(\hat{\theta} + \omega_k \Delta t) \\ \hat{\theta} + \omega_k \Delta t \\ \mathbf{m}_{k-1} \end{pmatrix}$$

Transition Update: Covariance

$$\mathbf{P}_{xx,k|k-1} = \nabla \mathbf{f} \mathbf{P}_{xx,k-1|k-1} \nabla \mathbf{f}^T + \mathbf{Q}_k$$

- As with the mean, covariance (uncertainty estimates) of the landmarks should not change when the robot moves

$$\nabla \mathbf{f} = \begin{pmatrix} \mathbf{F}_k & \mathbf{0} \\ \mathbf{0} & \mathbf{I} \end{pmatrix}$$

- Ex:

$$\mathbf{F}_k = \frac{\partial}{\partial (\hat{x}, \hat{y}, \hat{\theta})} \begin{pmatrix} \hat{x} - \frac{v_k}{\omega_k} \sin \hat{\theta} + \frac{v_k}{\omega_k} \sin(\hat{\theta} + \omega_k \Delta t) \\ \hat{y} + \frac{v_k}{\omega_k} \cos \hat{\theta} - \frac{v_k}{\omega_k} \cos(\hat{\theta} + \omega_k \Delta t) \\ \hat{\theta} + \omega_k \Delta t \end{pmatrix} = \begin{pmatrix} 1 & 0 & -\frac{v_k}{\omega_k} \cos \hat{\theta} + \frac{v_k}{\omega_k} \cos(\hat{\theta} + \omega_k \Delta t) \\ 0 & 1 & -\frac{v_k}{\omega_k} \sin \hat{\theta} + \frac{v_k}{\omega_k} \sin(\hat{\theta} + \omega_k \Delta t) \\ 0 & 0 & 1 \end{pmatrix}$$

Measurement Update

$$\mathbf{K}_k = \mathbf{P}_{k|k-1} \nabla \mathbf{h}^T (\nabla \mathbf{h} \mathbf{P}_{k|k-1} \nabla \mathbf{h}^T + \mathbf{R}_k)^{-1}$$

- Next step is to compute the Kalman to determine strength of updates
- We'll require the observation model here (or rather, its Jacobian)
- Example: Range-bearing measurement of landmark j


$$\mathbf{h}(\hat{\mathbf{x}}_{k|k-1}, \hat{\mathbf{m}}_{k-1}) = \begin{pmatrix} r_k \\ \phi_k \end{pmatrix} = \begin{pmatrix} \sqrt{(\hat{m}_{jx} - \hat{x})^2 + (\hat{m}_{jy} - \hat{y})^2} \\ \text{Atan2}(\hat{m}_{jy} - \hat{y}, \hat{m}_{jx} - \hat{x}) - \hat{\theta} \end{pmatrix}$$

- Jacobian components of all other landmarks should be 0!

Landmark Observations



Landmark Observations



Observation Jacobian

$$\mathbf{K}_k = \mathbf{P}_{k|k-1} \nabla \mathbf{h}^T (\nabla \mathbf{h} \mathbf{P}_{k|k-1} \nabla \mathbf{h}^T + \mathbf{R}_k)^{-1}$$

$$\mathbf{h}(\hat{\mathbf{x}}_{k|k-1}, \hat{\mathbf{m}}_{k-1}) = \begin{pmatrix} r_k \\ \phi_k \end{pmatrix} = \begin{pmatrix} \sqrt{(\hat{m}_{jx} - \hat{x})^2 + (\hat{m}_{jy} - \hat{y})^2} \\ \text{Atan2}(\hat{m}_{jy} - \hat{y}, \hat{m}_{jx} - \hat{x}) - \hat{\theta} \end{pmatrix}$$

$$\nabla \mathbf{h} = \frac{1}{r_k^2} \begin{pmatrix} -(\hat{m}_{jx} - \hat{x})r_k & -(\hat{m}_{jy} - \hat{y})r_k & 0 & 0 & \cdots & 0 & (\hat{m}_{jx} - \hat{x})r_k & (\hat{m}_{jy} - \hat{y})r_k & 0 & \cdots & 0 \\ (\hat{m}_{jy} - \hat{y}) & -(\hat{m}_{jx} - \hat{x}) & -r_k^2 & 0 & \cdots & 0 & -(\hat{m}_{jy} - \hat{y}) & (\hat{m}_{jx} - \hat{x}) & 0 & \cdots & 0 \end{pmatrix}$$

$\underbrace{\hspace{15em}}$
 Jacobian of \mathbf{h} wrt robot state $\hat{\mathbf{x}}$

$\underbrace{\hspace{5em}}$
 $2(j-1)$

$\underbrace{\hspace{15em}}$
 Jacobian of \mathbf{h} wrt j
 landmark, $\hat{m}_{jx}, \hat{m}_{jy}$

$\underbrace{\hspace{10em}}$
 $2(n-j)$

Adding New Landmarks

- What if we observe a new landmark instead of one that is already known?
- Invert the observation function and append to the state:

$$\mathbf{h}^{-1}(\hat{\mathbf{x}}_{k|k-1}, \mathbf{z}_k) = \begin{pmatrix} \hat{m}_{jx} \\ \hat{m}_{jy} \end{pmatrix} = \begin{pmatrix} \hat{x} + r_k \cos(\hat{\theta} + \phi_k) \\ \hat{y} + r_k \sin(\hat{\theta} + \phi_k) \end{pmatrix} \quad \hat{\mathbf{x}} \leftarrow \begin{pmatrix} \hat{x} \\ \hat{m}_{jx} \\ \hat{m}_{jy} \end{pmatrix}$$

- What about the covariance? In theory we can compute Jacobians of the observation function to obtain correlations between new landmark and current state covariances (robot state + landmarks)
- Simpler approach: Assign an estimate (e.g. observation covariance \mathbf{R}_k) with 0 cross-correlations, and covariances will hopefully converge to true values later on

EKF SLAM Algorithm

- Transition model: $\mathbf{x}_k = \mathbf{f}(\mathbf{x}_{k-1}, \mathbf{u}_k) + \mathbf{w}_k$ $\mathbf{w}_k \sim N(0, \mathbf{Q}_k)$
- Observation model: $\mathbf{z}_k = \mathbf{h}(\mathbf{x}_k, \mathbf{m}) + \mathbf{v}_k$ $\mathbf{v}_k \sim N(0, \mathbf{R}_k)$

- Given: $\hat{\mathbf{x}}_{k-1|k-1}, \mathbf{P}_{k-1|k-1}, \mathbf{u}_k, \mathbf{z}_k$

- Transition update: $\hat{\mathbf{x}}_{k|k-1} = \mathbf{f}(\hat{\mathbf{x}}_{k-1|k-1}, \mathbf{u}_k)$
 $\mathbf{P}_{xx,k|k-1} = \nabla \mathbf{f} \mathbf{P}_{xx,k-1|k-1} \nabla \mathbf{f}^T + \mathbf{Q}_k$


- Observation update: If landmark is new, add to state; otherwise:

$$\mathbf{K}_k = \mathbf{P}_{k|k-1} \nabla \mathbf{h}^T (\nabla \mathbf{h} \mathbf{P}_{k|k-1} \nabla \mathbf{h}^T + \mathbf{R}_k)^{-1}$$

$$\begin{pmatrix} \hat{\mathbf{x}}_{k|k} \\ \hat{\mathbf{m}}_k \end{pmatrix} = \begin{pmatrix} \hat{\mathbf{x}}_{k|k-1} \\ \hat{\mathbf{m}}_{k-1} \end{pmatrix} + \mathbf{K}_k (\mathbf{z}_k - \mathbf{h}(\hat{\mathbf{x}}_{k|k-1}, \hat{\mathbf{m}}_{k-1})) \quad \mathbf{P}_{k|k} = (\mathbf{I} - \mathbf{K}_k \nabla \mathbf{h}) \mathbf{P}_{k|k-1}$$


EKF SLAM Considerations

- Computational complexity increases quadratically in size of the map
- Highly nonlinear models may diverge or have suboptimal performance
- Convergence of EKF SLAM (if linear approximation is good): landmark estimates become fully correlated in the limit of infinitely many observations



Map

Covariance matrix





Map

Covariance matrix

Loop Closure

- **Loop closure:** Observing a landmark that has been observed already
- If the robot *believes* it has closed a loop, then uncertainties of landmarks it observed before will tend to collapse very quickly



Data Association Problem

- Success of EKF SLAM depends on robustness of data association, or ability to recognize an observed landmark (or a new one)
- Loop closure helps reduce uncertainties but can be disastrous if incorrect
- For best performance, landmarks should be distinct and spaced out
- Gives greater resilience to both motion and measurement noise
- If multiple possible candidates, maximum likelihood (ML) methods can help pick the best landmark, but most methods only rely on one such choice
- An incorrect choice due to ambiguity can lead to filter divergence!


Particle Filters for SLAM

- It would be better if we can maintain and update beliefs for all relevant landmarks simultaneously, maybe weighted by observation likelihoods
- A task well-suited for particle filters!

- Problem: How to represent distributions without combinatorial explosion?
- Naively, each particle would have to represent joint distribution between robot state and all landmarks


- Observation: Although all landmarks are correlated, we can consider them independent *conditioned on* robot's state!

Conditional Independence of Landmarks



Rao-Blackwellized Particle Filter

- Let's factor our joint belief distribution:

$$\begin{aligned} B(\mathbf{x}_{1:t}, \mathbf{m}_{1:M}) &= p(\mathbf{x}_{1:t}, \mathbf{m}_{1:M} | \mathbf{u}_{1:t}, \mathbf{z}_{1:t}) \\ &= p(\mathbf{x}_{1:t} | \mathbf{u}_{1:t}, \mathbf{z}_{1:t}) p(\mathbf{m}_{1:M} | \mathbf{x}_{1:t}, \mathbf{z}_{1:t}) \\ &= p(\mathbf{x}_{1:t} | \mathbf{u}_{1:t}, \mathbf{z}_{1:t}) \prod_{i=1}^M p(\mathbf{m}_i | \mathbf{x}_{1:t}, \mathbf{z}_{1:t}) \end{aligned}$$


State belief distribution can be maintained and updated using regular particle filter as before

Each particle maintains M EKFS, one for each landmark, and updates based on EKF process (we don't need to sample these!)

FastSLAM

- Each particle contains the following information
 - Robot state (assumed to be *correct* from perspective of landmarks)
 - Small (2 by 2), independent EKFs for each landmark conditioned on robot state
 - Particle weight

Particle 1

x, y, θ

Landmark 1

Landmark 2

... Landmark M

Particle 2

x, y, θ

Landmark 1

Landmark 2

... Landmark M

Particle N


x, y, θ

Landmark 1


Landmark 2

... Landmark M


FastSLAM Example: Transition Update



FastSLAM Example: Measurement Update



FastSLAM Example: Weights and Resampling



FastSLAM Algorithm

- At any given step, we start with a set of particles as described previously

- For each particle:

- Sample from transition model (move it): $p(\mathbf{x}_k | \mathbf{x}_{k-1}, \mathbf{u}_k)$

- Perform EKF observation update for observed landmark j :

$$\begin{aligned} \mathbf{K} &= \mathbf{P}_{j,k-1} \nabla \mathbf{h}^T (\nabla \mathbf{h} \mathbf{P}_{j,k-1} \nabla \mathbf{h}^T + \mathbf{R}_j)^{-1} & \hat{\mathbf{m}}_{j,k} &= \hat{\mathbf{m}}_{j,k-1} + \mathbf{K}(\mathbf{z}_k - \mathbf{h}(\hat{\mathbf{x}}_k, \hat{\mathbf{m}}_{j,k-1})) \\ \mathbf{P}_{j|k} &= (\mathbf{I} - \mathbf{K} \nabla \mathbf{h}) \mathbf{P}_{j,k-1} \end{aligned}$$

- Compute a corresponding weight (more on this below)

- Resample each particle according to current weights

Particle Weights

- With regular particle filters, particle weights were just observation likelihoods

$$p(\mathbf{z}_k | \mathbf{x}_k, \hat{\mathbf{m}})$$

- Since we have EKFs for landmarks, this likelihood can be computed analytically

- Covariance of **innovation**, or difference between true and predicted measurements, from the Kalman gain: $\Sigma_k = \nabla \mathbf{h} \mathbf{P}_{j,k-1} \nabla \mathbf{h}^T + \mathbf{R}_j$

- Particle weight (from Gaussian properties of EKF):

$$w_k = \frac{1}{\sqrt{2\pi|\Sigma_k|}} \exp\left(-\frac{1}{2}(\mathbf{z}_k - \hat{\mathbf{z}}_k)^T \Sigma_k^{-1}(\mathbf{z}_k - \hat{\mathbf{z}}_k)\right)$$

FastSLAM Considerations

- Conditional independence assumption allows us to sample only the robot's evolving state (nonlinear) while exploiting Gaussian nature of map beliefs
- As with regular particle filter, performance scales with number of particles
- Robust to data association ambiguities, can tolerate more nonlinearities
- Difficulties and problems of vanilla particle filters are still present
- Particle deprivation, over-dependence on measurement history, etc.