


COMS W4733: Computational Aspects of Robotics

Lecture 24: Computer Vision 1



Instructor: Tony Dear


Vision and Perception

- Any intelligent system requires some form of feedback from environment
- Vision is a specific form of perception based on images—easy to gather using cameras and many ways to process and interpret
- Relevant tasks: Camera calibration, image processing, object recognition, pose estimation, motion analysis, scene reconstruction, visual servoing
- Methods: Camera physics, image transformation, feature extraction, segmentation, closed-loop control

Image Processing




Semantic information



Geometric information


Segmentation

- Goal: Divide image up into meaningful components
 - Foreground vs background, different colors or shading, different classes of objects
- Related: Find significant features in image, e.g. edges or lines
- First problem: How to numerically represent images? Histograms!



- X-axis: Pixel value (typically, pure black is 0 and pure white is 255)
- Y-axis: Number of pixels with pixel value
- We discard a lot of information (e.g. spatial) this way!

Example: Image Histogram



Binary Segmentation

- While simple, image histogram can be used to perform simple **thresholding**
- Pixels belong to one of two labels, depending on value relative to a threshold
- Reveal locations of interesting features while compressing image

- Simple approach: Decide on a value μ
- All pixels greater than μ : make them 255; otherwise, make them 0
- How to pick μ ? Maybe start with 128

- More generally, use the histogram to decide on an appropriate threshold!


Choosing a Threshold

- The histogram is a “probability function” that a given pixel will have a given gray value
- One option: Use the image’s gray value mean as a threshold
- Good start, but can be improved since variance also provides information
 - Ex: Image with half white, half black vs image that is mostly gray
- Improved approach: Pick a threshold that would minimize some weighted combination of the two groups’ variances

Mean: $\mu_i = \sum_{z=0}^{255} \frac{H_i(z)}{\sum_{z=0}^{N-1} H_i(z)}$

Variance: $\sigma_i^2 = \sum_{z=0}^{255} (z - \mu_i)^2 \frac{H_i(z)}{\sum_{z=0}^{N-1} H_i(z)}$

Example: Thresholding



Connected Components

- Images often contain more than just foreground and background
- Labels should be further differentiated between different **components**
- Component identification can be done by simply noting adjacent pixels
- Raster scanning (left to right, up to down) image for connected components
- First pass—for each potential component pixel, assign to either existing component if adjacent to a seen one or assign to new one if not
 - If more than one adjacent component, mark down equivalence
- Second pass—reassign pixels after combining equivalences

Example: Raster Scanning Components


[illegible][illegible][illegible][illegible]

Example: Connected Components




Segmentation by Clustering

- Image segmentation can also be done via clustering methods
- Groups of “similar” pixels tend to clump together
- k -means allows one to specify number of components beforehand
- Randomly select k points as cluster means
- Repeat until convergence:
 - Assign each point to the closest mean
 - Move each mean to average of its assigned points





Boundary Segmentation

- Idea: Instead of grouping pixels, find local discontinuities in the grayness level
- We can define functions on pixels such that gradient on boundaries is large
- Edge detection: Take all pixels whose gradient is higher than some threshold
- Since images are discretized, “gradients” really just mean local differences



- Scanline of a row of pixels and gray values
- Gradient, or difference between values of successive pixels


Example: Edge Detection



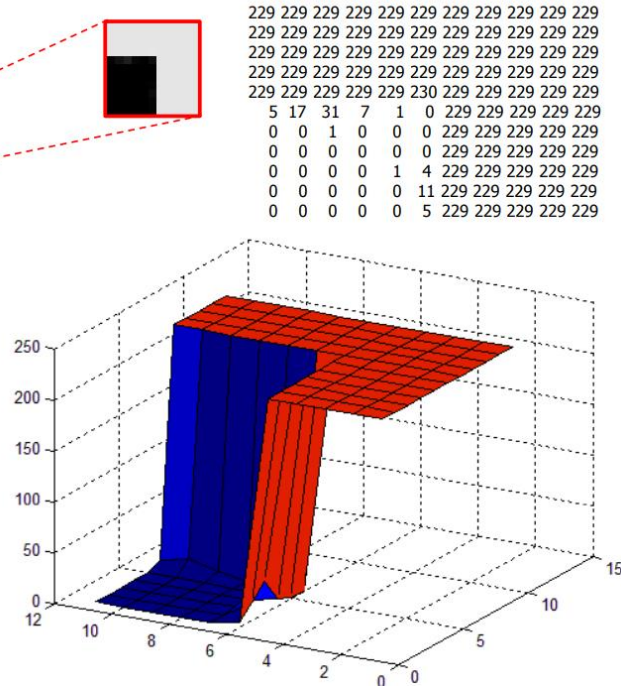
Example: Edge Detection



208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208
208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208
208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208
208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208
208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208
208	207	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208	208
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0



Example: Edge Detection



Edge Detection

- Process of taking local differences at each pixel is really a process of applying *masks* on the image pixels via convolution
- Horizontal edge: $[-1 \quad 1]^T$
- Vertical edge: $[-1 \quad 1]$
- More sophisticated masks also possible, e.g. Sobel operators


$$\text{Gradient magnitude: } \|G\| = \sqrt{G_x^2 + G_y^2}$$

$$\text{Horizontal: } \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

$$\text{Vertical: } \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

Masks and Smoothing

- Notion of applying masks to an image is important for another reason
- Real images are noisy! Taking differences on raw images may be a bad idea
- Smoothing should be applied first




Masks and Smoothing

- Convolution of a mask with an image: Replace each pixel value with dot product of mask and nearby pixel values
- Usually used for averaging (smoothing), but masks can be anything

- Ex: 3×3 mean averaging filter:
$$\begin{bmatrix} 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \end{bmatrix}$$

- Ex: Gaussian smoothing filter:




$$\frac{1}{273}$$


1	4	7	4	1
4	16	26	16	4
7	26	41	26	7
4	16	26	16	4
1	4	7	4	1

Example: After Smoothing


$$I(x)$$




$$G_{\sigma}(x)$$



$$s(x) = I(x) * G_{\sigma}(x)$$



$$s'(x) = \frac{d}{dx}(s(x))$$



Canny Edge Detector


- Combine smoothing with edge detection, plus some more steps for robustness
 - Criteria: Locate all edges, localize all edges, and minimize false edges
-
1. Gaussian smoothing to reduce noise
 2. Find intensity gradients (e.g. Sobel)
 3. Non-maximum suppression to *thin out* edges
 4. Double thresholding to identify strong edges and discard weak ones
 5. Edge tracking to see whether weak edges are isolated and should be removed

Feature Detection

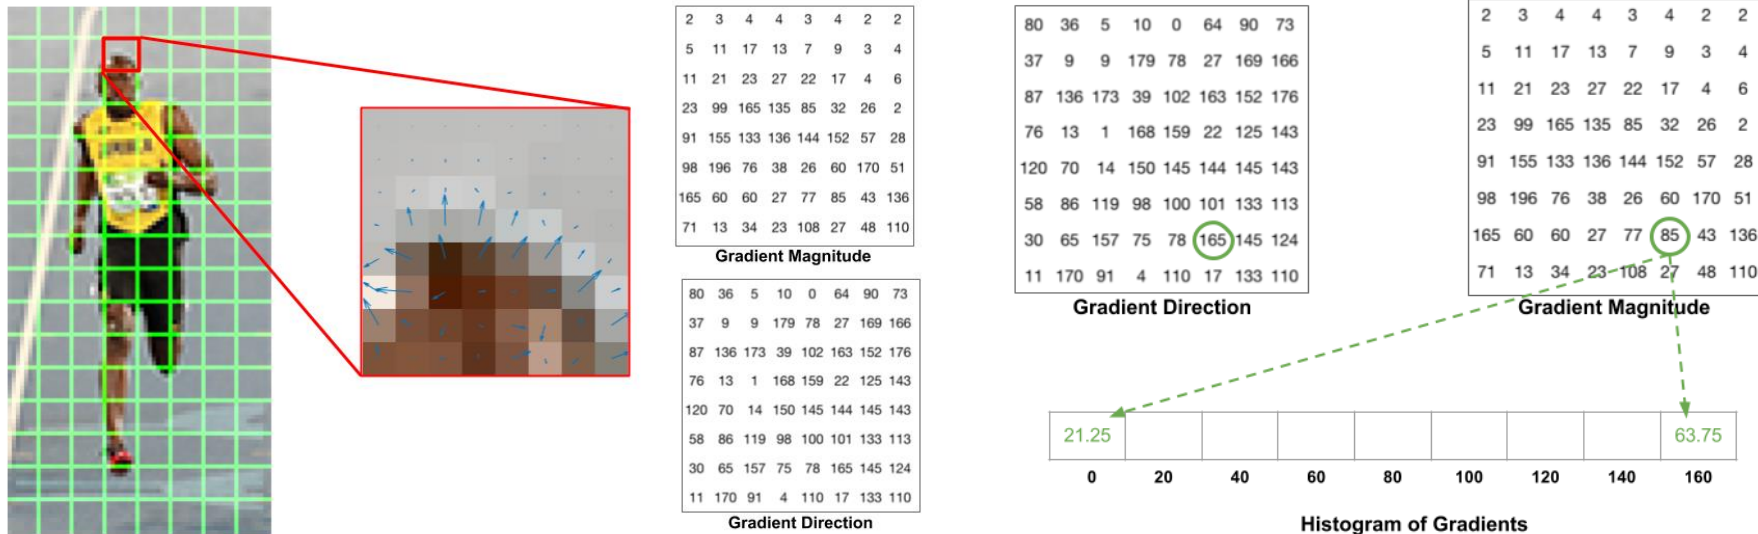
- In addition to being used for segmentation, edges can also serve as *features*
- Features: Any interesting part of an image useful for later tasks
- Should be repeatable (easily identifiable elsewhere on the image), may be examined at the pixel level
- Other feature types: Corners, blobs, ridges
- All have specialized detection algorithms
- Can we have some more generalized feature descriptors?

Histogram of Oriented Gradients

- Idea: Distribution of gradients in a local area may provide some useful information about object appearance and shape
- After computing gradient information at each pixel, look at local patches of discrete cells to find weighted histogram over all possible gradient angles
- All histograms in image are combined together as one giant feature vector





Example: HOGs



<https://www.learnopencv.com/histogram-of-oriented-gradients/>


Example: HOGs



Scale-Invariant Feature Transform



- We need to be able to match similar features across different images
- Same object might be viewed at different scales, angles, lighting, viewpoints

- High-level overview:



- Find key points as extrema on Gaussian-blurred image
- Refine localization of key points using thresholding
- Assign orientations to key points for robustness to rotations
- Create feature descriptor using orientation histograms

Example: SIFT



Summary

- Lots of different approaches in processing and understanding images!
- Segmentation: Dividing image up into meaningful components
- Feature detection: Find points of interest without extraneous information
- Techniques: Filters, thresholding, edge detection, HOGs, SIFT, ...
- Next time: Vision used in robotics, including pose estimation and visual control